



Vidyavardhini's College of Engineering & Technology

Department of Computer Engineering

Experiment No.4
Experiment on Hadoop Map-Reduce
Date of Performance: 14/08/23
Date of Submission: 21/08/23



Vidyavardhini's College of Engineering & Technology

Department of Computer Engineering

AIM: -To write a program to implement a word count program using MapReduce.

THEORY:

WordCount is a simple program which counts the number of occurrences of each word in a given text input data set. WordCount fits very well with the MapReduce programming model making it a great example to understand the Hadoop Map/Reduce programming style. The implementation consists of three main parts:

1. Mapper
2. Reducer
3. Driver

Step-1. Write a Mapper

A Mapper overrides the `map()` function from the Class "org.apache.hadoop.mapreduce.Mapper" which provides <key, value> pairs as the input. A Mapper implementation may output <key,value> pairs using the provided Context .

Input value of the WordCount Map task will be a line of text from the input data file and the key would be the line number <line_number, line_of_text> . Map task outputs <word, one> for each word in the line of text. Pseudo-code void Map (key, value){ for each word x in value:

```
output.collect(x,1);
}
```

Step-2. Write a Reducer

A Reducer collects the intermediate <key,value> output from multiple map tasks and assemble a single result. Here, the WordCount program will sum up the occurrence of each word to pairs as <word, occurrence>. Pseudo-code

```
void Reduce (keyword, <list of value>){ for each
x in <list of value>:
sum+=x;
final_output.collect(keyword, sum);
}
```



Vidyavardhini's College of Engineering & Technology

Department of Computer Engineering

Code:

```
import java.io.IOException; import
java.util.StringTokenizer; import
org.apache.hadoop.io.IntWritable; import
org.apache.hadoop.io.LongWritable; import
org.apache.hadoop.io.Text; import
org.apache.hadoop.mapreduce.Mapper; import
org.apache.hadoop.mapreduce.Reducer; import
org.apache.hadoop.conf.Configuration; import
org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.mapreduce.lib.input.TextInputFormat;
import org.apache.hadoop.mapreduce.lib.output.TextOutputFormat;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;
import org.apache.hadoop.fs.Path; public class WordCount
{
public static class Map extends Mapper<LongWritable,Text,Text,IntWritable> { public
void map(LongWritable key, Text value,Context context) throws
IOException,InterruptedException{
String line = value.toString();
StringTokenizer tokenizer = new StringTokenizer(line);
while      (tokenizer.hasMoreTokens())      {
value.set(tokenizer.nextToken());  context.write(value,
new IntWritable(1));
}
}
```

CSL702: Big Data Analytics Lab



Vidyavardhini's College of Engineering & Technology

Department of Computer Engineering

```
} }  
  
public static class Reduce extends Reducer<Text,IntWritable,Text,IntWritable> {  
    public void reduce(Text key, Iterable<IntWritable> values,Context context)  
        throws IOException,InterruptedException { int sum=0; for(IntWritable x: values)  
        { sum+=x.get();  
        }  
        context.write(key, new IntWritable(sum));  
    } }  
  
public static void main(String[] args) throws Exception {  
    Configuration conf= new Configuration(); Job job = new  
    Job(conf,"My Word Count Program");  
    job.setJarByClass(WordCount.class);  
    job.setMapperClass(Map.class);  
    job.setReducerClass(Reduce.class);  
    job.setOutputKeyClass(Text.class);  
    job.setOutputValueClass(IntWritable.class);  
    job.setInputFormatClass(TextInputFormat.class);  
    job.setOutputFormatClass(TextOutputFormat.class);  
  
    Path outputPath = new Path(args[1]);  
  
    //Configuring the input/output path from the filesystem into the job  
    FileInputFormat.addInputPath(job, new Path(args[0])); FileOutputFormat.setOutputPath(job,  
    new Path(args[1]));  
  
    //deleting the output path automatically from hdfs so that we don't have to delete  
    it explicitly  
    outputPath.getFileSystem(conf).delete(outputPath); //exiting  
    the job only if the flag value becomes false  
    CSL702: Big Data Analytics Lab
```



Vidyavardhini's College of Engineering & Technology

Department of Computer Engineering

```
System.exit(job.waitForCompletion(true) ? 0 : 1);
```

```
}
```

```
}
```

OUTPUT:



Vidyavardhini's College of Engineering & Technology

Department of Computer Engineering

```
package org.shubham;
import java.io.IOException;
import java.util.StringTokenizer;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.LongWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapred.MapReduceBase;
import org.apache.hadoop.mapred.Mapper;
import org.apache.hadoop.mapred.OutputCollector;
import org.apache.hadoop.mapred.Reporter;

1 usage
public class WC_Mapper extends MapReduceBase implements Mapper<LongWritable,Text,Text,IntWritable>{
    1 usage
    private final static IntWritable one = new IntWritable( value: 1);
    2 usages
    private Text word = new Text();
    public void map(LongWritable key, Text value,OutputCollector<Text,IntWritable> output,
        Reporter reporter) throws IOException{
        String line = value.toString();
```

```
package org.shubham;
import java.io.IOException;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapred.FileInputFormat;
import org.apache.hadoop.mapred.FileOutputFormat;
import org.apache.hadoop.mapred.JobClient;
import org.apache.hadoop.mapred.JobConf;
import org.apache.hadoop.mapred.TextInputFormat;
import org.apache.hadoop.mapred.TextOutputFormat;
public class WC_Runner {
    13 @
    public static void main(String[] args) throws IOException{
        JobConf conf = new JobConf(WC_Runner.class);
        conf.setJobName("WordCount");
        conf.setOutputKeyClass(Text.class);
        conf.setOutputValueClass(IntWritable.class);
        conf.setMapperClass(WC_Mapper.class);
        conf.setCombinerClass(WC_Reducer.class);
```

```
CONNECTION=0
IO_ERROR=0
WRONG_LENGTH=0
WRONG_MAP=0
WRONG_REDUCE=0

File Input Format Counters
  Bytes Read=32
File Output Format Counters
  Bytes Written=30

PS C:\Users\sadnak\IdeaProjects\WordCount>
```



Vidyavardhini's College of Engineering & Technology

Department of Computer Engineering

The screenshot displays the Hadoop Distributed File System (HDFS) Explorer interface. The browser address bar shows the URL `localhost:9870/explorer.html#/input`. The main content area is titled "Browse Directory" and shows the path `/input`. A modal window titled "File information - prac4.txt" is open, displaying the following details:

- Block information: Block 0
- Block ID: 1073741828
- Block Pool ID: BP-446742625-192.168.73.157-1697140882384
- Generation Stamp: 1004
- Size: 21
- Availability: Shubham

The "File contents" section shows the text: "Hello this is Shubham".

The browser address bar shows the URL `localhost:9870/explorer.html#/output`. The main content area is titled "Browse Directory" and shows the path `/output`. A modal window titled "File information - part-00000" is open, displaying the following details:

- Block information: Block 0
- Block ID: 1073741835
- Block Pool ID: BP-446742625-192.168.73.157-1697140882384
- Generation Stamp: 1011
- Size: 30
- Availability: Shubham

The "File contents" section shows the text: "Hello 1", "Shubham 1", "is 1", "this 1".



CONCLUSION:

The experiment aimed at implementing a word count program using MapReduce yielded successful results. MapReduce demonstrated its ability to scale and efficiently process extensive datasets by distributing tasks across multiple nodes for parallel execution. It also showcased its fault tolerance, ensuring the integrity of data processing in distributed systems. MapReduce's simplicity, characterized by its straightforward mapper and reducer functions, makes it accessible to a broad spectrum of developers. The practical relevance of this experiment extends to more intricate data processing tasks, such as log analysis and machine learning. By incorporating performance optimizations like combiners and partitioners, the program's efficiency can be further improved. In essence, this experiment established a robust foundation in distributed computing, a highly valuable skill in today's data-driven landscape.